

train_mobilenet_pneumonia_full

May 6, 2025

```
[1]: import os
import datetime
import torch
import torchvision.transforms as transforms
import torchvision.datasets as datasets
import torch.nn as nn
import torch.optim as optim
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from torchvision import models
from tqdm import tqdm
import sys
from sklearn.metrics import confusion_matrix, classification_report
from collections import Counter
from PIL import Image

[2]: DATA_DIR = os.path.expanduser("~/Zoidberg/data/raw/chest_Xray/test")
MODEL_PATH = os.path.expanduser("~/Zoidberg/models/mobilenet_pneumonia.pt")
DOCS_DIR = os.path.expanduser("~/Zoidberg/docs")
TIMESTAMP = datetime.datetime.now().strftime("%Y%m%d_%H%M%S")
BATCH_SIZE = 32
EPOCHS = 10
LEARNING_RATE = 1e-4
DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")

[3]: data_transforms = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
])

def load_dataset():
    class_map = {"NORMAL": 0, "bacterial": 1, "viral": 2}
    dataset = datasets.ImageFolder(DATA_DIR, transform=data_transforms)
    new_samples = []
    for path, label in dataset.samples:
        fname = os.path.basename(path).lower()
```

```

        if "pneumonia" in fname:
            label = class_map["bacterial"] if "bacteria" in fname else
↪class_map["viral"]
        else:
            label = class_map["NORMAL"]
            new_samples.append((path, label))
    dataset.samples = new_samples
    dataset.targets = [s[1] for s in new_samples]
    dataset.classes = ["NORMAL", "bacterial", "viral"]
    return dataset

dataset = load_dataset()
train_loader = torch.utils.data.DataLoader(dataset, batch_size=BATCH_SIZE,
↪shuffle=True)
val_loader = torch.utils.data.DataLoader(dataset, batch_size=BATCH_SIZE)

```

```

[4]: train_losses = []
    train_accuracies = []

def train():
    model = models.mobilenet_v2(weights=models.MobileNet_V2_Weights.DEFAULT)
    model.classifier[1] = nn.Linear(model.last_channel, 3)
    model = model.to(DEVICE)

    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters(), lr=LEARNING_RATE)

    for epoch in range(EPOCHS):
        model.train()
        running_loss = 0.0
        correct = 0
        total = 0
        loop = tqdm(train_loader, desc=f"Epoch {epoch+1}/{EPOCHS}", leave=True,
↪file=sys.stdout)
        for inputs, labels in loop:
            inputs, labels = inputs.to(DEVICE), labels.to(DEVICE)
            optimizer.zero_grad()
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()
            running_loss += loss.item()
            _, preds = torch.max(outputs, 1)
            correct += (preds == labels).sum().item()
            total += labels.size(0)
            loop.set_postfix(loss=running_loss / (total / BATCH_SIZE),
↪acc=correct / total)

```

```

        train_losses.append(running_loss / len(train_loader))
        train_accuracies.append(correct / total)
        print(f"Epoch {epoch+1} - Loss: {train_losses[-1]:.4f} - Accuracy:␣
↪{train_accuracies[-1]:.4f}")

    torch.save(model.state_dict(), MODEL_PATH)
    print(f"Modèle sauvegardé dans : {MODEL_PATH}")
    return model

model = train()

```

```

Epoch 1/10:  0%|                                     | 0/20 [00:00<?,
?it/s]

Epoch 1/10:  0%|                                     | 0/20 [00:06<?, ?it/s, acc=0.562,
loss=1.01]

Epoch 1/10:  5%|                                     | 1/20 [00:06<02:11,  6.93s/it, acc=0.562,
loss=1.01]

Epoch 1/10:  5%|                                     | 1/20 [00:10<02:11,  6.93s/it, acc=0.562,
loss=1.02]

Epoch 1/10: 10%|                                     | 2/20 [00:10<01:24,  4.67s/it, acc=0.562,
loss=1.02]

Epoch 1/10: 10%|                                     | 2/20 [00:12<01:24,  4.67s/it, acc=0.625,
loss=1]

Epoch 1/10: 15%|                                     | 3/20 [00:12<01:03,  3.72s/it, acc=0.625,
loss=1]

Epoch 1/10: 15%|                                     | 3/20 [00:15<01:03,  3.72s/it, acc=0.695,
loss=0.977]

Epoch 1/10: 20%|                                     | 4/20 [00:15<00:53,  3.37s/it, acc=0.695,
loss=0.977]

Epoch 1/10: 20%|                                     | 4/20 [00:17<00:53,  3.37s/it, acc=0.756,
loss=0.963]

Epoch 1/10: 25%|                                     | 5/20 [00:17<00:45,  3.04s/it, acc=0.756,
loss=0.963]

Epoch 1/10: 25%|                                     | 5/20 [00:20<00:45,  3.04s/it, acc=0.797,
loss=0.947]

Epoch 1/10: 30%|                                     | 6/20 [00:20<00:40,  2.88s/it, acc=0.797,
loss=0.947]

Epoch 1/10: 30%|                                     | 6/20 [00:23<00:40,  2.88s/it, acc=0.817,
loss=0.932]

Epoch 1/10: 35%|                                     | 7/20 [00:23<00:36,  2.84s/it, acc=0.817,
loss=0.932]

```

Epoch 1/10: 35%| | 7/20 [00:25<00:36, 2.84s/it, acc=0.84, loss=0.916]

Epoch 1/10: 40%| | 8/20 [00:25<00:31, 2.62s/it, acc=0.84, loss=0.916]

Epoch 1/10: 40%| | 8/20 [00:28<00:31, 2.62s/it, acc=0.858, loss=0.901]

Epoch 1/10: 45%| | 9/20 [00:28<00:29, 2.69s/it, acc=0.858, loss=0.901]

Epoch 1/10: 45%| | 9/20 [00:30<00:29, 2.69s/it, acc=0.872, loss=0.884]

Epoch 1/10: 50%| | 10/20 [00:30<00:26, 2.60s/it, acc=0.872, loss=0.884]

Epoch 1/10: 50%| | 10/20 [00:33<00:26, 2.60s/it, acc=0.884, loss=0.87]

Epoch 1/10: 55%| | 11/20 [00:33<00:23, 2.66s/it, acc=0.884, loss=0.87]

Epoch 1/10: 55%| | 11/20 [00:35<00:23, 2.66s/it, acc=0.893, loss=0.858]

Epoch 1/10: 60%| | 12/20 [00:35<00:20, 2.53s/it, acc=0.893, loss=0.858]

Epoch 1/10: 60%| | 12/20 [00:42<00:20, 2.53s/it, acc=0.901, loss=0.845]

Epoch 1/10: 65%| | 13/20 [00:42<00:26, 3.78s/it, acc=0.901, loss=0.845]

Epoch 1/10: 65%| | 13/20 [00:44<00:26, 3.78s/it, acc=0.908, loss=0.831]

Epoch 1/10: 70%| | 14/20 [00:44<00:19, 3.32s/it, acc=0.908, loss=0.831]

Epoch 1/10: 70%| | 14/20 [00:47<00:19, 3.32s/it, acc=0.915, loss=0.818]

Epoch 1/10: 75%| | 15/20 [00:47<00:15, 3.12s/it, acc=0.915, loss=0.818]

Epoch 1/10: 75%| | 15/20 [00:49<00:15, 3.12s/it, acc=0.92, loss=0.805]

Epoch 1/10: 80%| | 16/20 [00:49<00:11, 2.82s/it, acc=0.92, loss=0.805]

Epoch 1/10: 80%| | 16/20 [00:52<00:11, 2.82s/it, acc=0.925, loss=0.792]

Epoch 1/10: 85%| | 17/20 [00:52<00:08, 2.88s/it, acc=0.925, loss=0.792]

Epoch 1/10: 85%| | 17/20 [00:54<00:08, 2.88s/it, acc=0.929, loss=0.779]

Epoch 1/10: 90%| | 18/20 [00:54<00:05, 2.64s/it, acc=0.929, loss=0.779]

Epoch 1/10: 90%| | 18/20 [00:56<00:05, 2.64s/it, acc=0.933, loss=0.766]

Epoch 1/10: 95%| | 19/20 [00:56<00:02, 2.60s/it, acc=0.933, loss=0.766]

Epoch 1/10: 95%| | 19/20 [00:57<00:02, 2.60s/it, acc=0.934, loss=0.773]

Epoch 1/10: 100%| | 20/20 [00:57<00:00, 2.09s/it, acc=0.934, loss=0.773]

Epoch 1/10: 100%| | 20/20 [00:57<00:00, 2.89s/it, acc=0.934, loss=0.773]

Epoch 1 - Loss: 0.7536 - Accuracy: 0.9343

Epoch 2/10: 0%| | 0/20 [00:00<?, ?it/s]

Epoch 2/10: 0%| | 0/20 [00:02<?, ?it/s, acc=1, loss=0.475]

Epoch 2/10: 5%| | 1/20 [00:02<00:39, 2.08s/it, acc=1, loss=0.475]

Epoch 2/10: 5%| | 1/20 [00:04<00:39, 2.08s/it, acc=1, loss=0.466]

Epoch 2/10: 10%| | 2/20 [00:04<00:44, 2.46s/it, acc=1, loss=0.466]

Epoch 2/10: 10%| | 2/20 [00:06<00:44, 2.46s/it, acc=1, loss=0.458]

Epoch 2/10: 15%| | 3/20 [00:06<00:39, 2.32s/it, acc=1, loss=0.458]

Epoch 2/10: 15%| | 3/20 [00:09<00:39, 2.32s/it, acc=1, loss=0.448]

Epoch 2/10: 20%| | 4/20 [00:09<00:38, 2.41s/it, acc=1, loss=0.448]

Epoch 2/10: 20%| | 4/20 [00:11<00:38, 2.41s/it, acc=1, loss=0.437]

Epoch 2/10: 25%	5/20 [00:11<00:33, 2.27s/it, acc=1, loss=0.437]
Epoch 2/10: 25%	5/20 [00:15<00:33, 2.27s/it, acc=1, loss=0.428]
Epoch 2/10: 30%	6/20 [00:15<00:38, 2.78s/it, acc=1, loss=0.428]
Epoch 2/10: 30%	6/20 [00:19<00:38, 2.78s/it, acc=1, loss=0.42]
Epoch 2/10: 35%	7/20 [00:19<00:44, 3.41s/it, acc=1, loss=0.42]
Epoch 2/10: 35%	7/20 [00:21<00:44, 3.41s/it, acc=1, loss=0.411]
Epoch 2/10: 40%	8/20 [00:21<00:35, 2.94s/it, acc=1, loss=0.411]
Epoch 2/10: 40%	8/20 [00:24<00:35, 2.94s/it, acc=1, loss=0.405]
Epoch 2/10: 45%	9/20 [00:24<00:31, 2.88s/it, acc=1, loss=0.405]
Epoch 2/10: 45%	9/20 [00:26<00:31, 2.88s/it, acc=1, loss=0.397]
Epoch 2/10: 50%	10/20 [00:26<00:27, 2.70s/it, acc=1, loss=0.397]
Epoch 2/10: 50%	10/20 [00:29<00:27, 2.70s/it, acc=1, loss=0.389]
Epoch 2/10: 55%	11/20 [00:29<00:23, 2.59s/it, acc=1, loss=0.389]
Epoch 2/10: 55%	11/20 [00:31<00:23, 2.59s/it, acc=1, loss=0.382]
Epoch 2/10: 60%	12/20 [00:31<00:19, 2.40s/it, acc=1, loss=0.382]
Epoch 2/10: 60%	12/20 [00:33<00:19, 2.40s/it, acc=1, loss=0.375]
Epoch 2/10: 65%	13/20 [00:33<00:17, 2.44s/it, acc=1, loss=0.375]
Epoch 2/10: 65%	13/20 [00:36<00:17, 2.44s/it, acc=1, loss=0.368]
Epoch 2/10: 70%	14/20 [00:36<00:14, 2.40s/it, acc=1, loss=0.368]

Epoch 2/10: 70%| | 14/20 [00:38<00:14, 2.40s/it, acc=1, loss=0.362]

Epoch 2/10: 75%| | 15/20 [00:38<00:12, 2.48s/it, acc=1, loss=0.362]

Epoch 2/10: 75%| | 15/20 [00:40<00:12, 2.48s/it, acc=1, loss=0.355]

Epoch 2/10: 80%| | 16/20 [00:40<00:09, 2.39s/it, acc=1, loss=0.355]

Epoch 2/10: 80%| | 16/20 [00:43<00:09, 2.39s/it, acc=1, loss=0.349]

Epoch 2/10: 85%| | 17/20 [00:43<00:07, 2.47s/it, acc=1, loss=0.349]

Epoch 2/10: 85%| | 17/20 [00:45<00:07, 2.47s/it, acc=1, loss=0.342]

Epoch 2/10: 90%| | 18/20 [00:45<00:04, 2.40s/it, acc=1, loss=0.342]

Epoch 2/10: 90%| | 18/20 [00:48<00:04, 2.40s/it, acc=1, loss=0.335]

Epoch 2/10: 95%| | 19/20 [00:48<00:02, 2.48s/it, acc=1, loss=0.335]

Epoch 2/10: 95%| | 19/20 [00:53<00:02, 2.48s/it, acc=1, loss=0.338]

Epoch 2/10: 100%| | 20/20 [00:53<00:00, 3.18s/it, acc=1, loss=0.338]

Epoch 2/10: 100%| | 20/20 [00:53<00:00, 2.67s/it, acc=1, loss=0.338]

Epoch 2 - Loss: 0.3291 - Accuracy: 1.0000

Epoch 3/10: 0%| | 0/20 [00:00<?, ?it/s]

Epoch 3/10: 0%| | 0/20 [00:02<?, ?it/s, acc=1, loss=0.199]

Epoch 3/10: 5%| | 1/20 [00:02<00:44, 2.34s/it, acc=1, loss=0.199]

Epoch 3/10: 5%| | 1/20 [00:04<00:44, 2.34s/it, acc=1, loss=0.191]

Epoch 3/10: 10%| | 2/20 [00:04<00:43, 2.39s/it, acc=1, loss=0.191]

Epoch 3/10: 10%	2/20 [00:06<00:43, 2.39s/it, acc=1, loss=0.19]
Epoch 3/10: 15%	3/20 [00:06<00:38, 2.29s/it, acc=1, loss=0.19]
Epoch 3/10: 15%	3/20 [00:09<00:38, 2.29s/it, acc=1, loss=0.189]
Epoch 3/10: 20%	4/20 [00:09<00:36, 2.31s/it, acc=1, loss=0.189]
Epoch 3/10: 20%	4/20 [00:11<00:36, 2.31s/it, acc=1, loss=0.184]
Epoch 3/10: 25%	5/20 [00:11<00:34, 2.33s/it, acc=1, loss=0.184]
Epoch 3/10: 25%	5/20 [00:14<00:34, 2.33s/it, acc=1, loss=0.181]
Epoch 3/10: 30%	6/20 [00:14<00:33, 2.41s/it, acc=1, loss=0.181]
Epoch 3/10: 30%	6/20 [00:16<00:33, 2.41s/it, acc=1, loss=0.177]
Epoch 3/10: 35%	7/20 [00:16<00:30, 2.33s/it, acc=1, loss=0.177]
Epoch 3/10: 35%	7/20 [00:19<00:30, 2.33s/it, acc=1, loss=0.174]
Epoch 3/10: 40%	8/20 [00:19<00:29, 2.43s/it, acc=1, loss=0.174]
Epoch 3/10: 40%	8/20 [00:21<00:29, 2.43s/it, acc=1, loss=0.171]
Epoch 3/10: 45%	9/20 [00:21<00:25, 2.31s/it, acc=1, loss=0.171]
Epoch 3/10: 45%	9/20 [00:23<00:25, 2.31s/it, acc=1, loss=0.168]
Epoch 3/10: 50%	10/20 [00:23<00:23, 2.36s/it, acc=1, loss=0.168]
Epoch 3/10: 50%	10/20 [00:26<00:23, 2.36s/it, acc=1, loss=0.165]
Epoch 3/10: 55%	11/20 [00:26<00:21, 2.43s/it, acc=1, loss=0.165]
Epoch 3/10: 55%	11/20 [00:28<00:21, 2.43s/it, acc=1, loss=0.162]

Epoch 3/10: 60%| | 12/20 [00:28<00:18, 2.34s/it, acc=1, loss=0.162]

Epoch 3/10: 60%| | 12/20 [00:32<00:18, 2.34s/it, acc=1, loss=0.16]

Epoch 3/10: 65%| | 13/20 [00:32<00:19, 2.82s/it, acc=1, loss=0.16]

Epoch 3/10: 65%| | 13/20 [00:36<00:19, 2.82s/it, acc=1, loss=0.157]

Epoch 3/10: 70%| | 14/20 [00:36<00:19, 3.17s/it, acc=1, loss=0.157]

Epoch 3/10: 70%| | 14/20 [00:38<00:19, 3.17s/it, acc=1, loss=0.155]

Epoch 3/10: 75%| | 15/20 [00:38<00:14, 3.00s/it, acc=1, loss=0.155]

Epoch 3/10: 75%| | 15/20 [00:40<00:14, 3.00s/it, acc=1, loss=0.152]

Epoch 3/10: 80%| | 16/20 [00:40<00:10, 2.71s/it, acc=1, loss=0.152]

Epoch 3/10: 80%| | 16/20 [00:43<00:10, 2.71s/it, acc=1, loss=0.149]

Epoch 3/10: 85%| | 17/20 [00:43<00:07, 2.61s/it, acc=1, loss=0.149]

Epoch 3/10: 85%| | 17/20 [00:45<00:07, 2.61s/it, acc=1, loss=0.147]

Epoch 3/10: 90%| | 18/20 [00:45<00:05, 2.52s/it, acc=1, loss=0.147]

Epoch 3/10: 90%| | 18/20 [00:47<00:05, 2.52s/it, acc=1, loss=0.145]

Epoch 3/10: 95%| | 19/20 [00:47<00:02, 2.48s/it, acc=1, loss=0.145]

Epoch 3/10: 95%| | 19/20 [00:48<00:02, 2.48s/it, acc=1, loss=0.146]

Epoch 3/10: 100%| | 20/20 [00:48<00:00, 2.01s/it, acc=1, loss=0.146]

Epoch 3/10: 100%| | 20/20 [00:48<00:00, 2.44s/it, acc=1, loss=0.146]

Epoch 3 - Loss: 0.1421 - Accuracy: 1.0000

Epoch 4/10:	0%	0/20 [00:00<?,
?it/s]		
Epoch 4/10:	0%	0/20 [00:02<?, ?it/s, acc=1,
loss=0.0949]		
Epoch 4/10:	5%	1/20 [00:02<00:42, 2.22s/it, acc=1,
loss=0.0949]		
Epoch 4/10:	5%	1/20 [00:04<00:42, 2.22s/it, acc=1,
loss=0.091]		
Epoch 4/10:	10%	2/20 [00:04<00:40, 2.28s/it, acc=1,
loss=0.091]		
Epoch 4/10:	10%	2/20 [00:07<00:40, 2.28s/it, acc=1,
loss=0.0904]		
Epoch 4/10:	15%	3/20 [00:07<00:41, 2.44s/it, acc=1,
loss=0.0904]		
Epoch 4/10:	15%	3/20 [00:09<00:41, 2.44s/it, acc=1,
loss=0.088]		
Epoch 4/10:	20%	4/20 [00:09<00:38, 2.42s/it, acc=1,
loss=0.088]		
Epoch 4/10:	20%	4/20 [00:11<00:38, 2.42s/it, acc=1,
loss=0.0866]		
Epoch 4/10:	25%	5/20 [00:11<00:35, 2.40s/it, acc=1,
loss=0.0866]		
Epoch 4/10:	25%	5/20 [00:14<00:35, 2.40s/it, acc=1,
loss=0.0843]		
Epoch 4/10:	30%	6/20 [00:14<00:33, 2.40s/it, acc=1,
loss=0.0843]		
Epoch 4/10:	30%	6/20 [00:17<00:33, 2.40s/it, acc=1,
loss=0.0832]		
Epoch 4/10:	35%	7/20 [00:17<00:34, 2.64s/it, acc=1,
loss=0.0832]		
Epoch 4/10:	35%	7/20 [00:22<00:34, 2.64s/it, acc=1,
loss=0.0821]		
Epoch 4/10:	40%	8/20 [00:22<00:40, 3.34s/it, acc=1,
loss=0.0821]		
Epoch 4/10:	40%	8/20 [00:24<00:40, 3.34s/it, acc=1,
loss=0.0812]		
Epoch 4/10:	45%	9/20 [00:24<00:34, 3.14s/it, acc=1,
loss=0.0812]		

Epoch 4/10: 45%	9/20 [00:27<00:34, 3.14s/it, acc=1, loss=0.0803]
Epoch 4/10: 50%	10/20 [00:27<00:28, 2.83s/it, acc=1, loss=0.0803]
Epoch 4/10: 50%	10/20 [00:29<00:28, 2.83s/it, acc=1, loss=0.0792]
Epoch 4/10: 55%	11/20 [00:29<00:25, 2.79s/it, acc=1, loss=0.0792]
Epoch 4/10: 55%	11/20 [00:31<00:25, 2.79s/it, acc=1, loss=0.0777]
Epoch 4/10: 60%	12/20 [00:31<00:20, 2.58s/it, acc=1, loss=0.0777]
Epoch 4/10: 60%	12/20 [00:34<00:20, 2.58s/it, acc=1, loss=0.0765]
Epoch 4/10: 65%	13/20 [00:34<00:18, 2.67s/it, acc=1, loss=0.0765]
Epoch 4/10: 65%	13/20 [00:37<00:18, 2.67s/it, acc=1, loss=0.0757]
Epoch 4/10: 70%	14/20 [00:37<00:15, 2.55s/it, acc=1, loss=0.0757]
Epoch 4/10: 70%	14/20 [00:39<00:15, 2.55s/it, acc=1, loss=0.0744]
Epoch 4/10: 75%	15/20 [00:39<00:12, 2.57s/it, acc=1, loss=0.0744]
Epoch 4/10: 75%	15/20 [00:41<00:12, 2.57s/it, acc=1, loss=0.0734]
Epoch 4/10: 80%	16/20 [00:41<00:09, 2.46s/it, acc=1, loss=0.0734]
Epoch 4/10: 80%	16/20 [00:45<00:09, 2.46s/it, acc=1, loss=0.0725]
Epoch 4/10: 85%	17/20 [00:45<00:08, 2.73s/it, acc=1, loss=0.0725]
Epoch 4/10: 85%	17/20 [00:47<00:08, 2.73s/it, acc=1, loss=0.0719]
Epoch 4/10: 90%	18/20 [00:47<00:05, 2.59s/it, acc=1, loss=0.0719]
Epoch 4/10: 90%	18/20 [00:49<00:05, 2.59s/it, acc=1, loss=0.0711]

Epoch 4/10: 95%| | 19/20 [00:49<00:02, 2.54s/it, acc=1,
loss=0.0711]

Epoch 4/10: 95%| | 19/20 [00:50<00:02, 2.54s/it, acc=1,
loss=0.0719]

Epoch 4/10: 100%| | 20/20 [00:50<00:00, 2.03s/it, acc=1,
loss=0.0719]

Epoch 4/10: 100%| | 20/20 [00:50<00:00, 2.54s/it, acc=1,
loss=0.0719]

Epoch 4 - Loss: 0.0701 - Accuracy: 1.0000

Epoch 5/10: 0%| | 0/20 [00:00<?,
?it/s]

Epoch 5/10: 0%| | 0/20 [00:06<?, ?it/s, acc=1,
loss=0.05]

Epoch 5/10: 5%| | 1/20 [00:06<01:57, 6.16s/it, acc=1,
loss=0.05]

Epoch 5/10: 5%| | 1/20 [00:08<01:57, 6.16s/it, acc=1,
loss=0.0474]

Epoch 5/10: 10%| | 2/20 [00:08<01:13, 4.08s/it, acc=1,
loss=0.0474]

Epoch 5/10: 10%| | 2/20 [00:10<01:13, 4.08s/it, acc=1,
loss=0.0464]

Epoch 5/10: 15%| | 3/20 [00:10<00:53, 3.13s/it, acc=1,
loss=0.0464]

Epoch 5/10: 15%| | 3/20 [00:13<00:53, 3.13s/it, acc=1,
loss=0.0454]

Epoch 5/10: 20%| | 4/20 [00:13<00:47, 2.97s/it, acc=1,
loss=0.0454]

Epoch 5/10: 20%| | 4/20 [00:15<00:47, 2.97s/it, acc=1,
loss=0.0451]

Epoch 5/10: 25%| | 5/20 [00:15<00:41, 2.75s/it, acc=1,
loss=0.0451]

Epoch 5/10: 25%| | 5/20 [00:18<00:41, 2.75s/it, acc=1,
loss=0.0446]

Epoch 5/10: 30%| | 6/20 [00:18<00:38, 2.75s/it, acc=1,
loss=0.0446]

Epoch 5/10: 30%| | 6/20 [00:21<00:38, 2.75s/it, acc=1,
loss=0.0438]

Epoch 5/10: 35%	7/20 [00:21<00:35, 2.70s/it, acc=1, loss=0.0438]
Epoch 5/10: 35%	7/20 [00:23<00:35, 2.70s/it, acc=1, loss=0.0434]
Epoch 5/10: 40%	8/20 [00:23<00:31, 2.61s/it, acc=1, loss=0.0434]
Epoch 5/10: 40%	8/20 [00:25<00:31, 2.61s/it, acc=1, loss=0.0428]
Epoch 5/10: 45%	9/20 [00:25<00:27, 2.49s/it, acc=1, loss=0.0428]
Epoch 5/10: 45%	9/20 [00:28<00:27, 2.49s/it, acc=1, loss=0.0424]
Epoch 5/10: 50%	10/20 [00:28<00:25, 2.58s/it, acc=1, loss=0.0424]
Epoch 5/10: 50%	10/20 [00:30<00:25, 2.58s/it, acc=1, loss=0.0419]
Epoch 5/10: 55%	11/20 [00:30<00:21, 2.41s/it, acc=1, loss=0.0419]
Epoch 5/10: 55%	11/20 [00:33<00:21, 2.41s/it, acc=1, loss=0.0417]
Epoch 5/10: 60%	12/20 [00:33<00:20, 2.50s/it, acc=1, loss=0.0417]
Epoch 5/10: 60%	12/20 [00:36<00:20, 2.50s/it, acc=1, loss=0.0413]
Epoch 5/10: 65%	13/20 [00:36<00:18, 2.59s/it, acc=1, loss=0.0413]
Epoch 5/10: 65%	13/20 [00:42<00:18, 2.59s/it, acc=1, loss=0.0409]
Epoch 5/10: 70%	14/20 [00:42<00:21, 3.59s/it, acc=1, loss=0.0409]
Epoch 5/10: 70%	14/20 [00:44<00:21, 3.59s/it, acc=1, loss=0.0405]
Epoch 5/10: 75%	15/20 [00:44<00:15, 3.19s/it, acc=1, loss=0.0405]
Epoch 5/10: 75%	15/20 [00:46<00:15, 3.19s/it, acc=1, loss=0.04]
Epoch 5/10: 80%	16/20 [00:46<00:11, 2.99s/it, acc=1, loss=0.04]

Epoch 5/10: 80%| | 16/20 [00:49<00:11, 2.99s/it, acc=1,
loss=0.0395]

Epoch 5/10: 85%| | 17/20 [00:49<00:08, 2.80s/it, acc=1,
loss=0.0395]

Epoch 5/10: 85%| | 17/20 [00:52<00:08, 2.80s/it, acc=1,
loss=0.0391]

Epoch 5/10: 90%| | 18/20 [00:52<00:05, 2.86s/it, acc=1,
loss=0.0391]

Epoch 5/10: 90%| | 18/20 [00:54<00:05, 2.86s/it, acc=1,
loss=0.0386]

Epoch 5/10: 95%| | 19/20 [00:54<00:02, 2.62s/it, acc=1,
loss=0.0386]

Epoch 5/10: 95%| | 19/20 [00:55<00:02, 2.62s/it, acc=1,
loss=0.0391]

Epoch 5/10: 100%| | 20/20 [00:55<00:00, 2.08s/it, acc=1,
loss=0.0391]

Epoch 5/10: 100%| | 20/20 [00:55<00:00, 2.76s/it, acc=1,
loss=0.0391]

Epoch 5 - Loss: 0.0382 - Accuracy: 1.0000

Epoch 6/10: 0%| | 0/20 [00:00<?,
?it/s]

Epoch 6/10: 0%| | 0/20 [00:02<?, ?it/s, acc=1,
loss=0.0304]

Epoch 6/10: 5%| | 1/20 [00:02<00:48, 2.56s/it, acc=1,
loss=0.0304]

Epoch 6/10: 5%| | 1/20 [00:04<00:48, 2.56s/it, acc=1,
loss=0.0305]

Epoch 6/10: 10%| | 2/20 [00:04<00:42, 2.38s/it, acc=1,
loss=0.0305]

Epoch 6/10: 10%| | 2/20 [00:07<00:42, 2.38s/it, acc=1,
loss=0.0297]

Epoch 6/10: 15%| | 3/20 [00:07<00:41, 2.45s/it, acc=1,
loss=0.0297]

Epoch 6/10: 15%| | 3/20 [00:09<00:41, 2.45s/it, acc=1,
loss=0.0292]

Epoch 6/10: 20%| | 4/20 [00:09<00:37, 2.36s/it, acc=1,
loss=0.0292]

Epoch 6/10: 20%	4/20 [00:12<00:37, 2.36s/it, acc=1, loss=0.0284]
Epoch 6/10: 25%	5/20 [00:12<00:35, 2.39s/it, acc=1, loss=0.0284]
Epoch 6/10: 25%	5/20 [00:13<00:35, 2.39s/it, acc=1, loss=0.0281]
Epoch 6/10: 30%	6/20 [00:13<00:31, 2.25s/it, acc=1, loss=0.0281]
Epoch 6/10: 30%	6/20 [00:19<00:31, 2.25s/it, acc=1, loss=0.0282]
Epoch 6/10: 35%	7/20 [00:19<00:41, 3.21s/it, acc=1, loss=0.0282]
Epoch 6/10: 35%	7/20 [00:22<00:41, 3.21s/it, acc=1, loss=0.0279]
Epoch 6/10: 40%	8/20 [00:22<00:38, 3.23s/it, acc=1, loss=0.0279]
Epoch 6/10: 40%	8/20 [00:24<00:38, 3.23s/it, acc=1, loss=0.0278]
Epoch 6/10: 45%	9/20 [00:24<00:32, 2.97s/it, acc=1, loss=0.0278]
Epoch 6/10: 45%	9/20 [00:27<00:32, 2.97s/it, acc=1, loss=0.0275]
Epoch 6/10: 50%	10/20 [00:27<00:28, 2.81s/it, acc=1, loss=0.0275]
Epoch 6/10: 50%	10/20 [00:29<00:28, 2.81s/it, acc=1, loss=0.0272]
Epoch 6/10: 55%	11/20 [00:29<00:24, 2.67s/it, acc=1, loss=0.0272]
Epoch 6/10: 55%	11/20 [00:32<00:24, 2.67s/it, acc=1, loss=0.0269]
Epoch 6/10: 60%	12/20 [00:32<00:20, 2.61s/it, acc=1, loss=0.0269]
Epoch 6/10: 60%	12/20 [00:34<00:20, 2.61s/it, acc=1, loss=0.0267]
Epoch 6/10: 65%	13/20 [00:34<00:17, 2.53s/it, acc=1, loss=0.0267]
Epoch 6/10: 65%	13/20 [00:36<00:17, 2.53s/it, acc=1, loss=0.0265]

Epoch 6/10: 70%| | 14/20 [00:36<00:14, 2.46s/it, acc=1,
loss=0.0265]

Epoch 6/10: 70%| | 14/20 [00:39<00:14, 2.46s/it, acc=1,
loss=0.0264]

Epoch 6/10: 75%| | 15/20 [00:39<00:12, 2.45s/it, acc=1,
loss=0.0264]

Epoch 6/10: 75%| | 15/20 [00:41<00:12, 2.45s/it, acc=1,
loss=0.0262]

Epoch 6/10: 80%| | 16/20 [00:41<00:10, 2.51s/it, acc=1,
loss=0.0262]

Epoch 6/10: 80%| | 16/20 [00:44<00:10, 2.51s/it, acc=1,
loss=0.0259]

Epoch 6/10: 85%| | 17/20 [00:44<00:07, 2.42s/it, acc=1,
loss=0.0259]

Epoch 6/10: 85%| | 17/20 [00:46<00:07, 2.42s/it, acc=1,
loss=0.0256]

Epoch 6/10: 90%| | 18/20 [00:46<00:04, 2.43s/it, acc=1,
loss=0.0256]

Epoch 6/10: 90%| | 18/20 [00:48<00:04, 2.43s/it, acc=1,
loss=0.0254]

Epoch 6/10: 95%| | 19/20 [00:48<00:02, 2.37s/it, acc=1,
loss=0.0254]

Epoch 6/10: 95%| | 19/20 [00:50<00:02, 2.37s/it, acc=1,
loss=0.0258]

Epoch 6/10: 100%| | 20/20 [00:50<00:00, 2.08s/it, acc=1,
loss=0.0258]

Epoch 6/10: 100%| | 20/20 [00:50<00:00, 2.51s/it, acc=1,
loss=0.0258]

Epoch 6 - Loss: 0.0251 - Accuracy: 1.0000

Epoch 7/10: 0%| | 0/20 [00:00<?,
?it/s]

Epoch 7/10: 0%| | 0/20 [00:06<?, ?it/s, acc=1,
loss=0.0194]

Epoch 7/10: 5%| | 1/20 [00:06<01:58, 6.25s/it, acc=1,
loss=0.0194]

Epoch 7/10: 5%| | 1/20 [00:08<01:58, 6.25s/it, acc=1,
loss=0.0203]

Epoch 7/10: 10%	2/20 [00:08<01:09, 3.85s/it, acc=1, loss=0.0203]
Epoch 7/10: 10%	2/20 [00:10<01:09, 3.85s/it, acc=1, loss=0.0196]
Epoch 7/10: 15%	3/20 [00:10<00:53, 3.14s/it, acc=1, loss=0.0196]
Epoch 7/10: 15%	3/20 [00:12<00:53, 3.14s/it, acc=1, loss=0.0193]
Epoch 7/10: 20%	4/20 [00:12<00:43, 2.70s/it, acc=1, loss=0.0193]
Epoch 7/10: 20%	4/20 [00:15<00:43, 2.70s/it, acc=1, loss=0.0193]
Epoch 7/10: 25%	5/20 [00:15<00:42, 2.82s/it, acc=1, loss=0.0193]
Epoch 7/10: 25%	5/20 [00:17<00:42, 2.82s/it, acc=1, loss=0.0193]
Epoch 7/10: 30%	6/20 [00:17<00:36, 2.58s/it, acc=1, loss=0.0193]
Epoch 7/10: 30%	6/20 [00:20<00:36, 2.58s/it, acc=1, loss=0.0191]
Epoch 7/10: 35%	7/20 [00:20<00:33, 2.61s/it, acc=1, loss=0.0191]
Epoch 7/10: 35%	7/20 [00:22<00:33, 2.61s/it, acc=1, loss=0.019]
Epoch 7/10: 40%	8/20 [00:22<00:29, 2.45s/it, acc=1, loss=0.019]
Epoch 7/10: 40%	8/20 [00:25<00:29, 2.45s/it, acc=1, loss=0.019]
Epoch 7/10: 45%	9/20 [00:25<00:27, 2.48s/it, acc=1, loss=0.019]
Epoch 7/10: 45%	9/20 [00:27<00:27, 2.48s/it, acc=1, loss=0.0189]
Epoch 7/10: 50%	10/20 [00:27<00:23, 2.33s/it, acc=1, loss=0.0189]
Epoch 7/10: 50%	10/20 [00:29<00:23, 2.33s/it, acc=1, loss=0.0187]
Epoch 7/10: 55%	11/20 [00:29<00:21, 2.42s/it, acc=1, loss=0.0187]

Epoch 7/10: 55%	11/20 [00:32<00:21, 2.42s/it, acc=1, loss=0.0186]
Epoch 7/10: 60%	12/20 [00:32<00:18, 2.37s/it, acc=1, loss=0.0186]
Epoch 7/10: 60%	12/20 [00:34<00:18, 2.37s/it, acc=1, loss=0.0185]
Epoch 7/10: 65%	13/20 [00:34<00:16, 2.42s/it, acc=1, loss=0.0185]
Epoch 7/10: 65%	13/20 [00:40<00:16, 2.42s/it, acc=1, loss=0.0184]
Epoch 7/10: 70%	14/20 [00:40<00:21, 3.58s/it, acc=1, loss=0.0184]
Epoch 7/10: 70%	14/20 [00:42<00:21, 3.58s/it, acc=1, loss=0.0182]
Epoch 7/10: 75%	15/20 [00:42<00:15, 3.10s/it, acc=1, loss=0.0182]
Epoch 7/10: 75%	15/20 [00:45<00:15, 3.10s/it, acc=1, loss=0.018]
Epoch 7/10: 80%	16/20 [00:45<00:11, 2.96s/it, acc=1, loss=0.018]
Epoch 7/10: 80%	16/20 [00:47<00:11, 2.96s/it, acc=1, loss=0.0178]
Epoch 7/10: 85%	17/20 [00:47<00:08, 2.69s/it, acc=1, loss=0.0178]
Epoch 7/10: 85%	17/20 [00:50<00:08, 2.69s/it, acc=1, loss=0.0177]
Epoch 7/10: 90%	18/20 [00:50<00:05, 2.71s/it, acc=1, loss=0.0177]
Epoch 7/10: 90%	18/20 [00:52<00:05, 2.71s/it, acc=1, loss=0.0176]
Epoch 7/10: 95%	19/20 [00:52<00:02, 2.60s/it, acc=1, loss=0.0176]
Epoch 7/10: 95%	19/20 [00:53<00:02, 2.60s/it, acc=1, loss=0.0179]
Epoch 7/10: 100%	20/20 [00:53<00:00, 2.17s/it, acc=1, loss=0.0179]
Epoch 7/10: 100%	20/20 [00:53<00:00, 2.69s/it, acc=1, loss=0.0179]

Epoch 7 - Loss: 0.0174 - Accuracy: 1.0000

Epoch 8/10:	0%	0/20 [00:00<?,
?it/s]		
Epoch 8/10:	0%	0/20 [00:01<?, ?it/s, acc=1,
loss=0.0144]		
Epoch 8/10:	5%	1/20 [00:01<00:36, 1.90s/it, acc=1,
loss=0.0144]		
Epoch 8/10:	5%	1/20 [00:04<00:36, 1.90s/it, acc=1,
loss=0.0145]		
Epoch 8/10:	10%	2/20 [00:04<00:42, 2.36s/it, acc=1,
loss=0.0145]		
Epoch 8/10:	10%	2/20 [00:06<00:42, 2.36s/it, acc=1,
loss=0.0144]		
Epoch 8/10:	15%	3/20 [00:06<00:37, 2.23s/it, acc=1,
loss=0.0144]		
Epoch 8/10:	15%	3/20 [00:08<00:37, 2.23s/it, acc=1,
loss=0.0144]		
Epoch 8/10:	20%	4/20 [00:08<00:35, 2.22s/it, acc=1,
loss=0.0144]		
Epoch 8/10:	20%	4/20 [00:11<00:35, 2.22s/it, acc=1,
loss=0.0143]		
Epoch 8/10:	25%	5/20 [00:11<00:34, 2.31s/it, acc=1,
loss=0.0143]		
Epoch 8/10:	25%	5/20 [00:13<00:34, 2.31s/it, acc=1,
loss=0.0142]		
Epoch 8/10:	30%	6/20 [00:13<00:32, 2.29s/it, acc=1,
loss=0.0142]		
Epoch 8/10:	30%	6/20 [00:15<00:32, 2.29s/it, acc=1,
loss=0.0142]		
Epoch 8/10:	35%	7/20 [00:15<00:29, 2.30s/it, acc=1,
loss=0.0142]		
Epoch 8/10:	35%	7/20 [00:21<00:29, 2.30s/it, acc=1,
loss=0.0142]		
Epoch 8/10:	40%	8/20 [00:21<00:41, 3.50s/it, acc=1,
loss=0.0142]		
Epoch 8/10:	40%	8/20 [00:24<00:41, 3.50s/it, acc=1,
loss=0.0141]		

Epoch 8/10: 45%	9/20 [00:24<00:35, 3.22s/it, acc=1, loss=0.0141]
Epoch 8/10: 45%	9/20 [00:26<00:35, 3.22s/it, acc=1, loss=0.014]
Epoch 8/10: 50%	10/20 [00:26<00:28, 2.86s/it, acc=1, loss=0.014]
Epoch 8/10: 50%	10/20 [00:29<00:28, 2.86s/it, acc=1, loss=0.0139]
Epoch 8/10: 55%	11/20 [00:29<00:24, 2.74s/it, acc=1, loss=0.0139]
Epoch 8/10: 55%	11/20 [00:30<00:24, 2.74s/it, acc=1, loss=0.0137]
Epoch 8/10: 60%	12/20 [00:30<00:19, 2.47s/it, acc=1, loss=0.0137]
Epoch 8/10: 60%	12/20 [00:33<00:19, 2.47s/it, acc=1, loss=0.0137]
Epoch 8/10: 65%	13/20 [00:33<00:16, 2.42s/it, acc=1, loss=0.0137]
Epoch 8/10: 65%	13/20 [00:35<00:16, 2.42s/it, acc=1, loss=0.0137]
Epoch 8/10: 70%	14/20 [00:35<00:14, 2.42s/it, acc=1, loss=0.0137]
Epoch 8/10: 70%	14/20 [00:38<00:14, 2.42s/it, acc=1, loss=0.0136]
Epoch 8/10: 75%	15/20 [00:38<00:12, 2.52s/it, acc=1, loss=0.0136]
Epoch 8/10: 75%	15/20 [00:40<00:12, 2.52s/it, acc=1, loss=0.0135]
Epoch 8/10: 80%	16/20 [00:40<00:09, 2.46s/it, acc=1, loss=0.0135]
Epoch 8/10: 80%	16/20 [00:43<00:09, 2.46s/it, acc=1, loss=0.0135]
Epoch 8/10: 85%	17/20 [00:43<00:07, 2.44s/it, acc=1, loss=0.0135]
Epoch 8/10: 85%	17/20 [00:45<00:07, 2.44s/it, acc=1, loss=0.0133]
Epoch 8/10: 90%	18/20 [00:45<00:04, 2.36s/it, acc=1, loss=0.0133]

Epoch 8/10: 90%| | 18/20 [00:47<00:04, 2.36s/it, acc=1,
loss=0.0132]

Epoch 8/10: 95%| | 19/20 [00:47<00:02, 2.34s/it, acc=1,
loss=0.0132]

Epoch 8/10: 95%| | 19/20 [00:48<00:02, 2.34s/it, acc=1,
loss=0.0135]

Epoch 8/10: 100%| | 20/20 [00:48<00:00, 1.99s/it, acc=1,
loss=0.0135]

Epoch 8/10: 100%| | 20/20 [00:48<00:00, 2.44s/it, acc=1,
loss=0.0135]

Epoch 8 - Loss: 0.0132 - Accuracy: 1.0000

Epoch 9/10: 0%| | 0/20 [00:00<?,
?it/s]

Epoch 9/10: 0%| | 0/20 [00:01<?, ?it/s, acc=1,
loss=0.0127]

Epoch 9/10: 5%| | 1/20 [00:01<00:37, 1.97s/it, acc=1,
loss=0.0127]

Epoch 9/10: 5%| | 1/20 [00:08<00:37, 1.97s/it, acc=1,
loss=0.0114]

Epoch 9/10: 10%| | 2/20 [00:08<01:24, 4.70s/it, acc=1,
loss=0.0114]

Epoch 9/10: 10%| | 2/20 [00:10<01:24, 4.70s/it, acc=1,
loss=0.0114]

Epoch 9/10: 15%| | 3/20 [00:10<01:01, 3.64s/it, acc=1,
loss=0.0114]

Epoch 9/10: 15%| | 3/20 [00:12<01:01, 3.64s/it, acc=1,
loss=0.0113]

Epoch 9/10: 20%| | 4/20 [00:12<00:47, 2.97s/it, acc=1,
loss=0.0113]

Epoch 9/10: 20%| | 4/20 [00:15<00:47, 2.97s/it, acc=1,
loss=0.0112]

Epoch 9/10: 25%| | 5/20 [00:15<00:43, 2.91s/it, acc=1,
loss=0.0112]

Epoch 9/10: 25%| | 5/20 [00:18<00:43, 2.91s/it, acc=1,
loss=0.0111]

Epoch 9/10: 30%| | 6/20 [00:18<00:38, 2.74s/it, acc=1,
loss=0.0111]

Epoch 9/10: 30%	6/20 [00:20<00:38, 2.74s/it, acc=1, loss=0.011]
Epoch 9/10: 35%	7/20 [00:20<00:34, 2.62s/it, acc=1, loss=0.011]
Epoch 9/10: 35%	7/20 [00:22<00:34, 2.62s/it, acc=1, loss=0.0109]
Epoch 9/10: 40%	8/20 [00:22<00:29, 2.47s/it, acc=1, loss=0.0109]
Epoch 9/10: 40%	8/20 [00:25<00:29, 2.47s/it, acc=1, loss=0.0109]
Epoch 9/10: 45%	9/20 [00:25<00:26, 2.45s/it, acc=1, loss=0.0109]
Epoch 9/10: 45%	9/20 [00:27<00:26, 2.45s/it, acc=1, loss=0.0108]
Epoch 9/10: 50%	10/20 [00:27<00:24, 2.41s/it, acc=1, loss=0.0108]
Epoch 9/10: 50%	10/20 [00:29<00:24, 2.41s/it, acc=1, loss=0.0107]
Epoch 9/10: 55%	11/20 [00:29<00:21, 2.40s/it, acc=1, loss=0.0107]
Epoch 9/10: 55%	11/20 [00:32<00:21, 2.40s/it, acc=1, loss=0.0108]
Epoch 9/10: 60%	12/20 [00:32<00:18, 2.36s/it, acc=1, loss=0.0108]
Epoch 9/10: 60%	12/20 [00:34<00:18, 2.36s/it, acc=1, loss=0.0107]
Epoch 9/10: 65%	13/20 [00:34<00:16, 2.35s/it, acc=1, loss=0.0107]
Epoch 9/10: 65%	13/20 [00:36<00:16, 2.35s/it, acc=1, loss=0.0107]
Epoch 9/10: 70%	14/20 [00:36<00:13, 2.30s/it, acc=1, loss=0.0107]
Epoch 9/10: 70%	14/20 [00:42<00:13, 2.30s/it, acc=1, loss=0.0106]
Epoch 9/10: 75%	15/20 [00:42<00:17, 3.45s/it, acc=1, loss=0.0106]
Epoch 9/10: 75%	15/20 [00:44<00:17, 3.45s/it, acc=1, loss=0.0105]

Epoch 9/10: 80%| | 16/20 [00:44<00:12, 3.11s/it, acc=1,
loss=0.0105]

Epoch 9/10: 80%| | 16/20 [00:47<00:12, 3.11s/it, acc=1,
loss=0.0104]

Epoch 9/10: 85%| | 17/20 [00:47<00:08, 2.92s/it, acc=1,
loss=0.0104]

Epoch 9/10: 85%| | 17/20 [00:49<00:08, 2.92s/it, acc=1,
loss=0.0104]

Epoch 9/10: 90%| | 18/20 [00:49<00:05, 2.79s/it, acc=1,
loss=0.0104]

Epoch 9/10: 90%| | 18/20 [00:52<00:05, 2.79s/it, acc=1,
loss=0.0103]

Epoch 9/10: 95%| | 19/20 [00:52<00:02, 2.70s/it, acc=1,
loss=0.0103]

Epoch 9/10: 95%| | 19/20 [00:53<00:02, 2.70s/it, acc=1,
loss=0.0105]

Epoch 9/10: 100%| | 20/20 [00:53<00:00, 2.23s/it, acc=1,
loss=0.0105]

Epoch 9/10: 100%| | 20/20 [00:53<00:00, 2.68s/it, acc=1,
loss=0.0105]

Epoch 9 - Loss: 0.0103 - Accuracy: 1.0000

Epoch 10/10: 0%| | 0/20 [00:00<?,
?it/s]

Epoch 10/10: 0%| | 0/20 [00:01<?, ?it/s, acc=1,
loss=0.00893]

Epoch 10/10: 5%| | 1/20 [00:01<00:37, 1.98s/it, acc=1,
loss=0.00893]

Epoch 10/10: 5%| | 1/20 [00:04<00:37, 1.98s/it, acc=1,
loss=0.00909]

Epoch 10/10: 10%| | 2/20 [00:04<00:40, 2.26s/it, acc=1,
loss=0.00909]

Epoch 10/10: 10%| | 2/20 [00:06<00:40, 2.26s/it, acc=1,
loss=0.00911]

Epoch 10/10: 15%| | 3/20 [00:06<00:38, 2.28s/it, acc=1,
loss=0.00911]

Epoch 10/10: 15%| | 3/20 [00:09<00:38, 2.28s/it, acc=1,
loss=0.00901]

Epoch 10/10:	20%	4/20 [00:09<00:38, 2.42s/it, acc=1, loss=0.00901]
Epoch 10/10:	20%	4/20 [00:11<00:38, 2.42s/it, acc=1, loss=0.00907]
Epoch 10/10:	25%	5/20 [00:11<00:34, 2.33s/it, acc=1, loss=0.00907]
Epoch 10/10:	25%	5/20 [00:14<00:34, 2.33s/it, acc=1, loss=0.009]
Epoch 10/10:	30%	6/20 [00:14<00:34, 2.46s/it, acc=1, loss=0.009]
Epoch 10/10:	30%	6/20 [00:16<00:34, 2.46s/it, acc=1, loss=0.00883]
Epoch 10/10:	35%	7/20 [00:16<00:31, 2.40s/it, acc=1, loss=0.00883]
Epoch 10/10:	35%	7/20 [00:19<00:31, 2.40s/it, acc=1, loss=0.00881]
Epoch 10/10:	40%	8/20 [00:19<00:30, 2.51s/it, acc=1, loss=0.00881]
Epoch 10/10:	40%	8/20 [00:25<00:30, 2.51s/it, acc=1, loss=0.00879]
Epoch 10/10:	45%	9/20 [00:25<00:40, 3.65s/it, acc=1, loss=0.00879]
Epoch 10/10:	45%	9/20 [00:27<00:40, 3.65s/it, acc=1, loss=0.00869]
Epoch 10/10:	50%	10/20 [00:27<00:32, 3.25s/it, acc=1, loss=0.00869]
Epoch 10/10:	50%	10/20 [00:30<00:32, 3.25s/it, acc=1, loss=0.00866]
Epoch 10/10:	55%	11/20 [00:30<00:26, 2.97s/it, acc=1, loss=0.00866]
Epoch 10/10:	55%	11/20 [00:32<00:26, 2.97s/it, acc=1, loss=0.00862]
Epoch 10/10:	60%	12/20 [00:32<00:21, 2.69s/it, acc=1, loss=0.00862]
Epoch 10/10:	60%	12/20 [00:34<00:21, 2.69s/it, acc=1, loss=0.00856]
Epoch 10/10:	65%	13/20 [00:34<00:18, 2.69s/it, acc=1, loss=0.00856]


```

Epoch 10/10: 65%|          | 13/20 [00:36<00:18,  2.69s/it, acc=1,
loss=0.00844]

Epoch 10/10: 70%|          | 14/20 [00:36<00:14,  2.50s/it, acc=1,
loss=0.00844]

Epoch 10/10: 70%|          | 14/20 [00:39<00:14,  2.50s/it, acc=1,
loss=0.00837]

Epoch 10/10: 75%|          | 15/20 [00:39<00:12,  2.45s/it, acc=1,
loss=0.00837]

Epoch 10/10: 75%|          | 15/20 [00:41<00:12,  2.45s/it, acc=1,
loss=0.00831]

Epoch 10/10: 80%|          | 16/20 [00:41<00:09,  2.36s/it, acc=1,
loss=0.00831]

Epoch 10/10: 80%|          | 16/20 [00:43<00:09,  2.36s/it, acc=1,
loss=0.00827]

Epoch 10/10: 85%|          | 17/20 [00:43<00:07,  2.40s/it, acc=1,
loss=0.00827]

Epoch 10/10: 85%|          | 17/20 [00:46<00:07,  2.40s/it, acc=1,
loss=0.00821]

Epoch 10/10: 90%|          | 18/20 [00:46<00:04,  2.49s/it, acc=1,
loss=0.00821]

Epoch 10/10: 90%|          | 18/20 [00:49<00:04,  2.49s/it, acc=1,
loss=0.00817]

Epoch 10/10: 95%|          | 19/20 [00:49<00:02,  2.50s/it, acc=1,
loss=0.00817]

Epoch 10/10: 95%|          | 19/20 [00:49<00:02,  2.50s/it, acc=1,
loss=0.00838]

Epoch 10/10: 100%|         | 20/20 [00:49<00:00,  2.01s/it, acc=1,
loss=0.00838]

Epoch 10/10: 100%|         | 20/20 [00:49<00:00,  2.50s/it, acc=1,
loss=0.00838]

```

Epoch 10 - Loss: 0.0082 - Accuracy: 1.0000

Modèle sauvegardé dans : /home/zoidberg/Zoidberg/models/mobilenet_pneumonia.pt

```

[5]: model.eval()
     y_true, y_pred = [], []
     labels = [0, 1, 2]
     with torch.no_grad():
         for inputs, labels_batch in tqdm(val_loader, desc="Évaluation", file=sys.
           ↪stdout):

```

```

        outputs = model(inputs.to(DEVICE))
        _, preds = torch.max(outputs, 1)
        y_true.extend(labels_batch.numpy())
        y_pred.extend(preds.cpu().numpy())

print("Répartition réelle dans y_true :", Counter(y_true))
print("Répartition prédite dans y_pred :", Counter(y_pred))

cm = confusion_matrix(y_true, y_pred, labels=labels)
report = classification_report(y_true, y_pred, labels=labels,
    ↳target_names=dataset.classes, zero_division=0)
print(report)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=dataset.classes,
    ↳yticklabels=dataset.classes)
plt.title(f"Matrice de confusion ({TIMESTAMP})")
plt.xlabel("Prédit")
plt.ylabel("Réel")
plt.tight_layout()
plt.show()

```

Évaluation: 0% | 0/20 [00:00<?,
?it/s]

Évaluation: 5% | 1/20 [00:01<00:22,
1.17s/it]

Évaluation: 10% | 2/20 [00:02<00:21,
1.19s/it]

Évaluation: 15% | 3/20 [00:03<00:19,
1.16s/it]

Évaluation: 20% | 4/20 [00:07<00:36,
2.28s/it]

Évaluation: 25% | 5/20 [00:09<00:32,
2.20s/it]

Évaluation: 30% | 6/20 [00:10<00:24,
1.74s/it]

Évaluation: 35% | 7/20 [00:11<00:19,
1.47s/it]

Évaluation: 40% | 8/20 [00:12<00:15,
1.33s/it]

Évaluation: 45% | 9/20 [00:13<00:13,
1.18s/it]

Évaluation: 50% | 10/20 [00:13<00:09,
1.01it/s]

Évaluation: 55%	11/20 [00:14<00:07,
1.16it/s]	
Évaluation: 60%	12/20 [00:14<00:06,
1.29it/s]	
Évaluation: 65%	13/20 [00:15<00:05,
1.27it/s]	
Évaluation: 70%	14/20 [00:16<00:04,
1.21it/s]	
Évaluation: 75%	15/20 [00:17<00:04,
1.23it/s]	
Évaluation: 80%	16/20 [00:18<00:03,
1.19it/s]	
Évaluation: 85%	17/20 [00:18<00:02,
1.28it/s]	
Évaluation: 90%	18/20 [00:19<00:01,
1.38it/s]	
Évaluation: 95%	19/20 [00:20<00:00,
1.45it/s]	
Évaluation: 100%	20/20 [00:20<00:00,
1.77it/s]	
Évaluation: 100%	20/20 [00:20<00:00,
1.02s/it]	

Répartition réelle dans y_true : Counter({np.int64(0): 624})

Répartition prédite dans y_pred : Counter({np.int64(0): 624})

	precision	recall	f1-score	support
NORMAL	1.00	1.00	1.00	624
bacterial	0.00	0.00	0.00	0
viral	0.00	0.00	0.00	0
accuracy			1.00	624
macro avg	0.33	0.33	0.33	624
weighted avg	1.00	1.00	1.00	624

